

Basics

Fundamental Assumption

Data is iid for unknown  $P$ :  $(x_i, y_i) \sim P(X, Y)$

True risk and estimated error

True risk:  $R(w) = \int P(x, y)(y - w^T x)^2 \partial x \partial y = \mathbb{E}_{x, y}[(y - w^T x)^2]$

Est. error:  $\hat{R}_D(w) = \frac{1}{|D|} \sum_{(x, y) \in D} (y - w^T x)^2$

Standardization

Centered data with unit variance:  $\tilde{x}_i = \frac{x_i - \hat{\mu}}{\hat{\sigma}}$   
 $\hat{\mu} = \frac{1}{n} \sum_{i=1}^n x_i, \hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \hat{\mu})^2$

Cross-Validation

For all models  $m$ , for all  $i \in \{1, \dots, k\}$  do:

1. Split data:  $D = D_{train}^{(i)} \uplus D_{test}^{(i)}$  (Monte-Carlo or k-Fold)
  2. Train model:  $\hat{w}_{i,m} = \underset{w}{\operatorname{argmin}} \hat{R}_{train}^{(i)}(w)$
  3. Estimate error:  $\hat{R}_m^{(i)} = \hat{R}_{test}^{(i)}(\hat{w}_{i,m})$
- Select best model:  $\hat{m} = \underset{m}{\operatorname{argmin}} \frac{1}{k} \sum_{i=1}^k \hat{R}_m^{(i)}$

Parametric vs. Nonparametric models

*Parametric*: have finite set of parameters. e.g. linear regression, linear perceptron

*Nonparametric*: grow in complexity with the size of the data, more expressive. e.g. k-NN

Gradient Descent

1. Pick arbitrary  $w_0 \in \mathbb{R}^d$
2.  $w_{t+1} = w_t - \eta_t \nabla \hat{R}(w_t)$

Stochastic Gradient Descent (SGD)

1. Pick arbitrary  $w_0 \in \mathbb{R}^d$
2.  $w_{t+1} = w_t - \eta_t \nabla_w l(w_t; x', y')$ , with u.a.r. data point  $(x', y') \in D$

Regression

Solve  $w^* = \underset{w}{\operatorname{argmin}} \hat{R}(w) + \lambda C(w)$

Linear Regression

$\hat{R}(w) = \sum_{i=1}^n (y_i - w^T x_i)^2 = \|Xw - y\|_2^2$   
 $\nabla_w \hat{R}(w) = -2 \sum_{i=1}^n (y_i - w^T x_i) \cdot x_i$   
 $w^* = (X^T X)^{-1} X^T y$

Ridge regression

$\hat{R}(w) = \sum_{i=1}^n (y_i - w^T x_i)^2 + \lambda \|w\|_2^2$   
 $\nabla_w \hat{R}(w) = -2 \sum_{i=1}^n (y_i - w^T x_i) \cdot x_i + 2\lambda w$   
 $w^* = (X^T X + \lambda I)^{-1} X^T y$

L1-regularized regression (Lasso)

$\hat{R}(w) = \sum_{i=1}^n (y_i - w^T x_i)^2 + \lambda \|w\|_1$

Classification

Solve  $w^* = \underset{w}{\operatorname{argmin}} l(w; x_i, y_i)$ ; loss function  $l$

0/1 loss

$l_{0/1}(w; y_i, x_i) = 1$  if  $y_i \neq \operatorname{sign}(w^T x_i)$  else 0

Perceptron algorithm

Use  $l_P(w; y_i, x_i) = \max(0, -y_i w^T x_i)$  and SGD

$$\nabla_w l_P(w; y_i, x_i) = \begin{cases} 0 & \text{if } y_i w^T x_i \geq 0 \\ -y_i x_i & \text{otherwise} \end{cases}$$

Data lin. separable  $\Leftrightarrow$  obtains a lin. separator (not necessarily optimal)

Support Vector Machine (SVM)

Hinge loss:  $l_H(w; x_i, y_i) = \max(0, 1 - y_i w^T x_i)$

$$\nabla_w l_H(w; y, x) = \begin{cases} 0 & \text{if } y_i w^T x_i \geq 1 \\ -y_i x_i & \text{otherwise} \end{cases}$$

$w^* = \underset{w}{\operatorname{argmin}} l_H(w; x_i, y_i) + \lambda \|w\|_2^2$

Kernels

efficient, implicit inner products

Properties of kernel

$k: \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ ,  $k$  must be some inner product (symmetric, positive-definite, linear) for some space  $\mathcal{V}$ . i.e.  $k(\mathbf{x}, \mathbf{x}') = \langle \varphi(\mathbf{x}), \varphi(\mathbf{x}') \rangle_{\mathcal{V}} \stackrel{\text{Eucl.}}{=} \varphi(\mathbf{x})^T \varphi(\mathbf{x}')$  and  $k(\mathbf{x}, \mathbf{x}') = k(\mathbf{x}', \mathbf{x})$

Kernel matrix

$$K = \begin{bmatrix} k(x_1, x_1) & \dots & k(x_1, x_n) \\ \vdots & \ddots & \vdots \\ k(x_n, x_1) & \dots & k(x_n, x_n) \end{bmatrix}$$

Positive semi-definite matrices  $\Leftrightarrow$  kernels  $k$

Important kernels

Linear:  $k(x, y) = x^T y$   
Polynomial:  $k(x, y) = (x^T y + 1)^d$   
Gaussian:  $k(x, y) = \exp(-\|x - y\|_2^2 / (2h^2))$   
Laplacian:  $k(x, y) = \exp(-\|x - y\|_1 / h)$

Composition rules

Valid kernels  $k_1, k_2$ , also valid kernels:  $k_1 + k_2$ ;  $k_1 \cdot k_2$ ;  $c \cdot k_1$ ,  $c > 0$ ;  $f(k_1)$  if  $f$  polynomial with pos. coeffs. or exponential

Reformulating the perceptron

Ansatz:  $w^* \in \operatorname{span}(X) \Rightarrow w = \sum_{j=1}^n \alpha_j y_j x_j$   
 $\alpha^* = \underset{\alpha \in \mathbb{R}^n}{\operatorname{argmin}} \sum_{i=1}^n \max(0, -\sum_{j=1}^n \alpha_j y_i y_j x_i^T x_j)$

Kernelized perceptron and SVM

Use  $\alpha^T k_i$  instead of  $w^T x_i$ ,  
use  $\alpha^T D_y K D_y \alpha$  instead of  $\|w\|_2^2$   
 $k_i = [y_1 k(x_i, x_1), \dots, y_n k(x_i, x_n)]$ ,  $D_y = \operatorname{diag}(y)$   
Prediction:  $\hat{y} = \operatorname{sign}(\sum_{i=1}^n \alpha_i y_i k(x_i, \hat{x}))$   
SGD update:  $\alpha_{t+1} = \alpha_t$ , if mispredicted:  
 $\alpha_{t+1, i} = \alpha_{t, i} + \eta_t$  (c.f. updating weights towards mispredicted point)

Kernelized linear regression (KLR)

Ansatz:  $w^* = \sum_{i=1}^n \alpha_i x$   
 $\alpha^* = \underset{\alpha}{\operatorname{argmin}} \|\alpha^T K - y\|_2^2 + \lambda \alpha^T K \alpha$   
 $= (K + \lambda I)^{-1} y$

Prediction:  $\hat{y} = \sum_{i=1}^n \alpha_i k(x_i, \hat{x})$

k-NN

$y = \operatorname{sign}(\sum_{i=1}^n y_i [x_i \text{ among } k \text{ nearest neighbours of } x])$  – No weights  $\Rightarrow$  no training! But depends on all data : (

Imbalance

Cost Sensitive Classification

Replace loss by:  $l_{CS}(w; x, y) = c_y l(w; x, y)$

Metrics

$n = n_+ + n_-$ ,  $n_+ = TP + FN$ ,  $n_- = TN + FP$   
Accuracy:  $\frac{TP + TN}{n}$ , Precision:  $\frac{TP}{TP + FP}$   
Recall/TPR:  $\frac{TP}{n_+}$ , FPR:  $\frac{FP}{n_-}$   
F1 score:  $\frac{2TP}{2TP + FP + FN}$   
ROC Curve:  $y = \text{TPR}$ ,  $x = \text{FPR}$

Multi-class

Hinge loss

$l_{MC-H}(w^{(1)}, \dots, w^{(c)}; x, y) = \max(0, 1 + \max_{j \in \{1, \dots, y-1, y+1, \dots, c\}} w^{(j)T} x - w^{(y)T} x)$

Neural network

Parameterize feature map:  $\phi(x, \theta)$  instead of  $\phi(x)$ , usually:  $\phi(x, \theta) = \varphi(\theta^T x) = \varphi(z)$   
 $\Rightarrow w^* = \underset{w, \theta}{\operatorname{argmin}} \sum_{i=1}^n l(y_i; \sum_{j=1}^m w_j \phi(x_i, \theta_j))$

Activation functions

Sigmoid:  $\frac{1}{1 + \exp(-z)}$ ,  $\varphi'(z) = (1 - \varphi(z)) \cdot \varphi(z)$   
Tanh:  $\varphi(z) = \tanh(z) = \frac{\exp(z) - \exp(-z)}{\exp(z) + \exp(-z)}$   
ReLU:  $\varphi(z) = \max(z, 0)$

Predict: forward propagation

$v^{(0)} = x$ ; for  $l = 1, \dots, L-1$ :  
 $v^{(l)} = \varphi(z^{(l)})$ ,  $z^{(l)} = W^{(l)} v^{(l-1)}$   
 $f = W^{(L)} v^{(L-1)}$   
Predict  $f$  for regression,  $\operatorname{sign}(f)$  for class.

Compute gradient: backpropagation

Output layer:  $\delta_j = l'_j(f_j)$ ,  $\frac{\partial}{\partial w_{j,i}} = \delta_j v_i$   
Hidden layer  $l = L-1, \dots, 1$ :  
 $\delta_j = \varphi'(z_j) \cdot \sum_{i \in \text{Layer}_{l+1}} w_{i,j} \delta_i$ ,  $\frac{\partial}{\partial w_{j,i}} = \delta_j v_i$

Learning with momentum

$a \leftarrow m \cdot a + \eta_t \nabla_W l(W; y, x)$ ;  $W \leftarrow W - a$

Clustering

k-mean

$\hat{R}(\mu) = \sum_{i=1}^n \min_{j \in \{1, \dots, k\}} \|x_i - \mu_j\|_2^2$   
 $\hat{\mu} = \underset{\mu}{\operatorname{argmin}} \hat{R}(\mu)$  ...non-convex, NP-hard

Algorithm (Lloyd's heuristic): Choose starting centers, assign points to closest center, update centers to mean of each cluster, repeat

Dimension reduction

PCA

$D = x_1, \dots, x_n \subset \mathbb{R}^d$ ,  $\Sigma = \frac{1}{n} \sum_{i=1}^n x_i x_i^T$ ,  $\mu = 0$   
 $(W, z_1, \dots, z_n) = \underset{W}{\operatorname{argmin}} \sum_{i=1}^n \|W z_i - x_i\|_2^2$ ,  
 $W = (v_1 | \dots | v_k) \in \mathbb{R}^{d \times k}$ , orthogonal;  $z_i = W^T x_i$   
 $v_i$  are the eigen vectors of  $\Sigma$

Kernel PCA

Kernel PC:  $\alpha^{(1)}, \dots, \alpha^{(k)} \in \mathbb{R}^n$ ,  $\alpha^{(i)} = \frac{1}{\sqrt{\lambda_i}} v_i$ ,  
 $K = \sum_{i=1}^n \lambda_i v_i v_i^T$ ,  $\lambda_1 \geq \dots \geq \lambda_d \geq 0$   
New point:  $\hat{z} = f(\hat{x}) = \sum_{j=1}^n \alpha_j^{(i)} k(\hat{x}, x_j)$

Autoencoders

Find identity function:  $x \approx f(x; \theta)$   
 $f(x; \theta) = f_{\text{decode}}(f_{\text{encode}}(x; \theta_{\text{encode}}); \theta_{\text{decode}})$

Probability modeling

Find  $h: X \rightarrow Y$  that min. pred. error:  $R(h) = \int P(x,y)l(y;h(x))\partial y x \partial y = \mathbb{E}_{x,y}[l(y;h(x))]$

For least squares regression

Best  $h$ :  $h^*(x) = \mathbb{E}[Y|X=x]$   
Pred.:  $\hat{y} = \hat{\mathbb{E}}[Y|X=\hat{x}] = \int \hat{P}(y|X=\hat{x})y\partial y$

Maximum Likelihood Estimation (MLE)

$\theta^* = \underset{\theta}{\operatorname{argmax}} \hat{P}(y_1,...,y_n|x_1,...,x_n,\theta)$   
E.g. lin. + Gauss:  $y_i = w^T x_i + \varepsilon_i, \varepsilon_i \sim \mathcal{N}(0,\sigma^2)$   
i.e.  $y_i \sim \mathcal{N}(w^T x_i, \sigma^2)$ , With MLE (use  $\operatorname{argmin} -\log$ ):  $w^* = \underset{w}{\operatorname{argmin}} \sum (y_i - w^T x_i)^2$

Bias/Variance/Noise

Prediction error =  $Bias^2 + Variance + Noise$

Maximum a posteriori estimate (MAP)

Assume bias on parameters, e.g.  $w_i \in \mathcal{N}(0,\beta^2)$   
Bay.:  $P(w|x,y) = \frac{P(w)P(y|x,w)}{P(y|x)} = \frac{P(w)P(y|x,w)}{P(y|x)}$

Logistic regression

Link func.:  $\sigma(w^T x) = \frac{1}{1+exp(-w^T x)}$  (Sigmoid)  
 $P(y|x,w) = Ber(y;\sigma(w^T x)) = \frac{1}{1+exp(-yw^T x)}$   
Classification: Use  $P(y|x,w)$ , predict most likely class label.  
MLE:  $\underset{w}{\operatorname{argmax}} P(y_{1:n}|w,x_{1:n})$   
 $\Rightarrow w^* = \underset{w}{\operatorname{argmin}} \sum_{i=1}^n \log(1+exp(-y_i w^T x_i))$

SGD update:  $w = w + \eta_i y_i \hat{P}(Y = -y|w,x)$   
 $\hat{P}(Y = -y|w,x) = \frac{1}{1+exp(yw^T x)}$   
MAP: Gauss. prior  $\Rightarrow ||w||_2^2$ , Lap. p.  $\Rightarrow ||w||_1$   
SGD:  $w = w(1 - 2\lambda \eta_t) + \eta_t y x \hat{P}(Y = -y|w,x)$

Bayesian decision theory

- Conditional distribution over labels  $P(y|x)$   
- Set of actions  $\mathcal{A}$   
- Cost function  $C: Y \times \mathcal{A} \rightarrow \mathbb{R}$   
 $a^* = \underset{a \in \mathcal{A}}{\operatorname{argmin}} \mathbb{E}[C(y,a)|x]$   
Calculate  $\mathbb{E}$  via sum/integral.  
Classification:  $C(y,a) = [y \neq a]$ ; asymmetric:  
$$C(y,a) = \begin{cases} c_{FP} & , \text{ if } y = -1, a = +1 \\ c_{FN} & , \text{ if } y = +1, a = -1 \\ 0 & , \text{ otherwise} \end{cases}$$
  
Regression:  $C(y,a) = (y - a)^2$ ; asymmetric:

$C(y,a) = c_1 \max(y - a, 0) + c_2 \max(a - y, 0)$   
E.g.  $y \in \{-1, +1\}$ , predict  $+$  if  $c_+ < c_-$ ,  
 $c_+ = \mathbb{E}(C(y, +1)|x) = P(y = 1|x) \cdot 0 + P(y = -1|x) \cdot c_{FP}$ ,  $c_-$  likewise

Discriminative / generative modeling

Discr. estimate  $P(y|x)$ , generative  $P(y,x)$   
Approach (generative):  $P(x,y) = P(x|y) \cdot P(y)$   
- Estimate prior on labels  $P(y)$   
- Estimate cond. distr.  $P(x|y)$  for each class  $y$   
- Pred. using Bayes:  $P(y|x) = \frac{P(y)P(x|y)}{P(x)}$   
 $P(x) = \sum_y P(x,y)$

Examples

MLE for  $P(y) = p = \frac{n_+}{n}$   
MLE for  $P(x_i|y) = \mathcal{N}(x_i; \mu_{i,y}, \sigma_{i,y}^2)$ :  
 $\hat{\mu}_{i,y} = \frac{1}{n_y} \sum_{x \in D_{x_i|y}} x$   
 $\hat{\sigma}_{i,y}^2 = \frac{1}{n_y} \sum_{x \in D_{x_i|y}} (x - \hat{\mu}_{i,y})^2$   
MLE for Poi.:  $\lambda = \operatorname{avg}(x_i)$   
 $\mathbb{R}^d$ :  $P(X = x|Y = y) = \prod_{i=1}^d Pois(\lambda_y^{(i)}, x^{(i)})$

Deriving decision rule

$P(y|x) = \frac{1}{Z} P(y)P(x|y)$ ,  $Z = \sum_y P(y)P(x|y)$   
 $y^* = \underset{y}{\operatorname{amax}} P(y|x) = \underset{y}{\operatorname{amax}} P(y) \prod_{i=1}^d P(x_i|y)$

Gaussian Bayes Classifier

$\hat{P}(x|y) = \mathcal{N}(x; \hat{\mu}_y, \hat{\Sigma}_y)$   
 $\hat{P}(Y = y) = \hat{p}_y = \frac{n_y}{n}$   
 $\hat{\mu}_y = \frac{1}{n_y} \sum_{i: y_i = y} x_i \in \mathbb{R}^d$   
 $\hat{\Sigma}_y = \frac{1}{n_y} \sum_{i: y_i = y} (x_i - \hat{\mu}_y)(x_i - \hat{\mu}_y)^T \in \mathbb{R}^{d \times d}$

Fisher's lin. discrim. analysis (LDA, c=2)

Assume:  $p = 0.5$ ;  $\hat{\Sigma}_- = \hat{\Sigma}_+ = \hat{\Sigma}$   
discriminant function:  $f(x) = \log \frac{p}{1-p} + \frac{1}{2} [\log \frac{|\hat{\Sigma}_-|}{|\hat{\Sigma}_+|} + ((x - \hat{\mu}_-)^T \hat{\Sigma}_-^{-1} (x - \hat{\mu}_-) - ((x - \hat{\mu}_+)^T \hat{\Sigma}_+^{-1} (x - \hat{\mu}_+))]$   
Predict:  $y = \operatorname{sign}(f(x)) = \operatorname{sign}(w^T x + w_0)$   
 $w = \hat{\Sigma}^{-1} (\hat{\mu}_+ - \hat{\mu}_-)$ ;  
 $w_0 = \frac{1}{2} (\hat{\mu}_-^T \hat{\Sigma}^{-1} \hat{\mu}_- - \hat{\mu}_+^T \hat{\Sigma}^{-1} \hat{\mu}_+)$

Outlier Detection

$P(x) \leq \tau$

Categorical Naive Bayes Classifier

MLE for feature distr.:  $\hat{P}(X_i = c|Y = y) = \theta_{c|y}^{(i)}$   
 $\theta_{c|y}^{(i)} = \frac{Count(X_i = c, Y = y)}{Count(Y = y)}$   
Prediction:  $y^* = \underset{y}{\operatorname{argmax}} \hat{P}(y|x)$

Missing data

Mixture modeling

Model each c. as probability distr.  $P(x|\theta_j)$   
 $P(D|\theta) = \prod_{i=1}^n \sum_{j=1}^k w_j P(x_i|\theta_j)$   
 $L(w, \theta) = -\sum_{i=1}^n \log \sum_{j=1}^k w_j P(x_i|\theta_j)$

Gaussian-Mixture Bayes classifiers

Estimate prior  $P(y)$ ; Est. cond. distr. for each class:  $P(x|y) = \sum_{j=1}^{k_y} w_j^{(y)} \mathcal{N}(x; \mu_j^{(y)}, \Sigma_j^{(y)})$

Hard-EM algorithm

Initialize parameters  $\theta^{(0)}$   
E-step: Predict most likely class for each point:  $z_i^{(t)} = \underset{z}{\operatorname{argmax}} P(z|x_i, \theta^{(t-1)})$   
 $= \underset{z}{\operatorname{argmax}} P(z|\theta^{(t-1)}) P(x_i|z, \theta^{(t-1)})$ ;  
M-step: Compute the MLE:  $\theta^{(t)} = \underset{\theta}{\operatorname{argmax}} P(D^{(t)}|\theta)$ , i.e.  $\mu_j^{(t)} = \frac{1}{n_j} \sum_{i: z_i = j} x_j$

Soft-EM algorithm

E-step: Calc p for each point and cls.:  $\gamma_j^{(t)}(x_i)$   
M-step: Fit clusters to weighted data points:  
 $w_j^{(t)} = \frac{1}{n} \sum_{i=1}^n \gamma_j^{(t)}(x_i)$ ;  $\mu_j^{(t)} = \frac{\sum_{i=1}^n \gamma_j^{(t)}(x_i) x_i}{\sum_{i=1}^n \gamma_j^{(t)}(x_i)}$   
 $\sigma_j^{(t)} = \frac{\sum_{i=1}^n \gamma_j^{(t)}(x_i) (x_i - \mu_j^{(t)})^T (x_i - \mu_j^{(t)})}{\sum_{i=1}^n \gamma_j^{(t)}(x_i)}$

Soft-EM for semi-supervised learning

labeled  $y_i$ :  $\gamma_j^{(t)}(x_i) = [j = y_i]$ , unlabeled:  
 $\gamma_j^{(t)}(x_i) = P(Z = j|x_i, \mu^{(t-1)}, \Sigma^{(t-1)}, w^{(t-1)})$

Useful math

Probabilities

$\mathbb{E}_x[X] = \begin{cases} \int x \cdot p(x) \partial x & \text{if continuous} \\ \sum_x x \cdot p(x) & \text{otherwise} \end{cases}$   
 $\operatorname{Var}[X] = \mathbb{E}[(X - \mu_X)^2] = \mathbb{E}[X^2] - \mathbb{E}[X]^2$

$P(A|B) = \frac{P(B|A) \cdot P(A)}{P(B)}$ ;  $p(Z|X, \theta) = \frac{p(X, Z|\theta)}{p(X|\theta)}$   
 $P(x,y) = P(y|x) \cdot P(x) = P(x|y) \cdot P(y)$

Bayes Rule

$P(A|B) = \frac{P(B|A) \cdot P(A)}{P(B)}$

P-Norm

$||x||_p = (\sum_{i=1}^n |x_i|^p)^{\frac{1}{p}}$ ,  $1 \leq p < \infty$   
Some gradients  
 $\nabla_x ||x||_2^2 = 2x$   
 $f(x) = x^T A x$ ;  $\nabla_x f(x) = (A + A^T)x$   
E.g.  $\nabla_w \log(1 + \exp(-yw^T x)) = \frac{1}{1 + \exp(-yw^T x)} \cdot \exp(-yw^T x) \cdot (-yx) = \frac{1}{1 + \exp(yw^T x)} \cdot (-yx)$

Convex / Jensen's inequality

$g(x)$  convex  $\Leftrightarrow g''(x) > 0 \Leftrightarrow x_1, x_2 \in \mathbb{R}, \lambda \in [0, 1]$ :  
 $g(\lambda x_1 + (1 - \lambda)x_2) \leq \lambda g(x_1) + (1 - \lambda)g(x_2)$

Gaussian / Normal Distribution

$f(x) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp(-\frac{(x-\mu)^2}{2\sigma^2})$

Multivariate Gaussian

$\Sigma$  = covariance matrix,  $\mu$  = mean  
 $f(x) = \frac{1}{2\pi\sqrt{|\Sigma|}} e^{-\frac{1}{2}(x-\mu)^T \Sigma^{-1}(x-\mu)}$   
Empirical:  $\hat{\Sigma} = \frac{1}{n} \sum_{i=1}^n x_i x_i^T$  (needs centered data points)

Positive semi-definite matrices

$M \in \mathbb{R}^{n \times n}$  is psd  $\Leftrightarrow \forall x \in \mathbb{R}^n : x^T M x \geq 0 \Leftrightarrow$   
all eigenvalues of  $M$  are positive:  $\lambda_i \geq 0$